

Question of the day

- How can we organize data to efficiently locate and update the most important elements?
- How can we efficiently store and update group memberships?

Priority queues and Union-Find

William Hendrix

Lecture 5


Outline

- Review
 - AVL trees
 - Hash tables
 - Sets
 - Maps
- Priority queues
- Heaps
- Union-Find


AVL tree review

- Self-balancing binary search tree
- All ops guaranteed $O(\lg n)$
- Nodes keep track of *balance*
 - $\text{height}(\text{right}) - \text{height}(\text{left})$
 - Always 0 or ± 1
- Performs *rotations* as nodes are inserted or deleted to maintain balance
 - 4 cases: left-left, left-right, right-left, right-right

Left rotation



Right rotation



Hash table review

- Sparse array-based structure
 - Values inserted based on *hash function*
- **Separate chaining:** array of linked lists
- **Open addressing:** insert at "next available" space
 - Quadratic probing and double hashing reduce congestion
- Rehash once *load factor* exceeds threshold
 - **Load factor:** size / capacity
 - Double capacity (roughly) and reinsert using hash function
 - Often chosen to be prime
- Hash functions
 - Need to be fast, distribute values evenly, and separate nearby values
 - Often use multiplication, polynomials, or bitwise ops
 - Mod at end to ensure range is 0 to *cap* - 1
 - *Multi-byte inputs:* multiply or rotate before incorporating next values

Set review

- ADT for storing and retrieving values
 - May allow or disallow duplicates
- **Main operations:**
 - *Search(x)*: returns whether x is in the set
 - *Insert(x)*: adds x to the set (may or may not allow duplicates)
 - *Delete(x)*: removes x from the set
- **Two main implementations:** balanced BST and hash table
 - Hash table "usually" faster
 - $\Theta(1)$ expected complexity
 - Assumes $\Theta(1)$ collisions
 - BST has better worst-case complexity
 - $O(\lg n)$ vs. $O(n)$
 - BST can access elements in sorted order
 - `min()`, `max()`, `predecessor()`, `successor()`

Maps review

- Stores set of associations
- **Main operations:**
 - *Insert(key, value)*: associates *value* to *key*
 - *Delete(key)*: removes any association with *key*
 - *Search(key)*: returns value associated with *key*
- **Implementations**
 - Array map:
 - Store value in `arr[key]`
 - $\Theta(1)$ worst case (*very fast!*)
 - Keys must be relatively small ints
 - Hash map:
 - Insert (key, value) pairs into hash table
 - Only hash key
 - $\Theta(1)$ expected complexity
 - Tree map:
 - Insert (key, value) pairs into BBST based on key
 - $O(\lg n)$ worst case

Maps

- Abstraction of a function
- Main operations
 - **Insert(x, y):** declares that $f(x) = y$
 - **Delete(x):** declares that $f(x)$ does not have a value
 - **Search(x):** returns y such that $f(x) = y$, or NIL if $f(x)$ does not have a value
- **Example:** letter frequencies
 - Problem: count how many times a letter appears in a given text
 - Used in cryptography
 - Sample output

E	T	A	O	I	N	S	R	H	...
12	9	8	7	7	6	6	6	6	...


- Need to associate count with letter
- $f(E) = 12$, etc.

Map implementation

- Two main implementations
- Array-based map
 - Array of all possible x values
 - Stores $f(x)$ in $\text{arr}[x]$ (NIL if not initialized)

8	2	3	4	12	2	2	...
A	B	C	D	E	F	G	...

- All main operations are constant time
 - Only useful when input domain is small
- Set-based map
 - A.k.a., hash map
 - Set of ordered (x, y) pairs
 - Pairs added/searched according to x value
 - **Search** returns associated y value
 - Time complexity determined by hash table or BBST



Map complexity

Operation	Array-based map	Set-based (hash table)	Set-based (BBST)
Insert(x, y)	$\Theta(1)$	$\Theta(1)^*$	$O(\lg n)$
Delete(x)	$\Theta(1)$	$\Theta(1)^*$	$O(\lg n)$
Search(x)	$\Theta(1)$	$\Theta(1)^*$	$O(\lg n)$
Build()	$\Theta(D)$	$\Theta(n)$	$O(n \lg n)$

D : size of domain (x values)

* Expected complexity for hash table

The power of maps

- Maps are very useful for storing values that we compute repeatedly
 - Especially when we can use direct maps
- **Example:** Discrete Fourier Transform
 - Given array x compute transformed array c such that

$$c_k = \sum_{j=1}^n x_j e^{jk(-2\pi i/n)} \longrightarrow \text{Store values in lookup table}$$


- Can also improve best-case performance for any algorithm
 1. Build a map that contains problem instances and solutions
 2. Before running another algorithm, test whether input is in map
 3. If so, return the answer
- Best case typically constant or linear time
- Best case analysis not useful to compare algorithm quality

Priority queues

- Abstract data type
- Similar to a queue, but returns elements in *priority order*
 - Min-first or max-first
 - Very good at finding min/max
- 3 main operations (min)
- **Insert(x)**
 - Adds another element
- **Min()**
 - Returns min
- **DeleteMin()**
 - Returns min and deletes


Heaps

- Main implementation of priority queue
- Complete binary tree that satisfies *heap property*
 - Min heap: every parent is \leq its children
 - Max heap: every parent is \geq its children
 - Unlike BST, left and right children not related
 - Also, children fill last level left-to-right
- **Min():**
 - Return root




Heap insertion

- **Insert(x):**
- Add x into next position on bottom level
 - Tree size dictates where to go
 - Increment size
 - Convert size to binary
 - Skip to just past first 1
 - Go left on 0
 - Go right on 1
- Restore heap property



Heap insertion

- **Insert(x):**
- Add x into next position on bottom level
 - Tree size dictates where to go
 - Increment size
 - Convert size to binary
 - Skip to just past first 1
 - Go left on 0
 - Go right on 1
- Restore heap property




Size = 6

110

Heap insertion

- **Insert(x):**
- Add x into next position on bottom level
 - Tree size dictates where to go
 - Increment size
 - Convert size to binary
 - Skip to just past first 1
 - Go left on 0
 - Go right on 1
- Restore heap property




Size = 6

110

Heap insertion

- **Insert(x):**
- Add x into next position on bottom level
 - Tree size dictates where to go
 - Increment size
 - Convert size to binary
 - Skip to just past first 1
 - Go left on 0
 - Go right on 1
- Restore heap property
 - Swap with parent if out-of-order
 - Repeat until satisfied or at root
 - "Percolate up"

*Bitwise
ops not
tested*




Size = 7

111


Heap example

- Add 35 to heap at right
- Then add 6




Heap example

- Add 35 to heap at right
- Then add 6




Heap exercise

- Sketch the result of inserting 6, 5, 4, 3, 2, and 1 into an empty min-heap




Heap exercise

- Sketch the result of inserting 6, 5, 4, 3, 2, and 1 into an empty min-heap




Heap deletion

- **DeleteMin():**
- Swap root value with last node
- Delete last node




Heap deletion

- **DeleteMin():**
- Swap root value with last node
- Delete last node
- Fix heap property at root
 - Swap root with min child
 - Max child for max-heap
 - Stop if node greater than both children or at leaf
 - "Percolate down"



Heap deletion

- **DeleteMin():**
- Swap root value with last node
- Delete last node
- Fix heap property at root
 - Swap root with min child
 - Max child for max-heap
 - Stop if node greater than both children or at leaf
 - "Percolate down"





Heap complexity

- **Min():**
 - Return root
 - $O(1)$
- **Insert(x):**
 - Scan to bottom of tree
 - Percolate up
 - Worst case: go back to root
 - $O(h) = O(\lg n)$
- **DeleteMin():**
 - Scan to bottom
 - Swap
 - Percolate down
 - Worst case: go down to leaf
 - $O(\lg n)$



Array-based heap

- Preferred implementation for a heap
- Store node values in an array
- Root at index 1
- Children of index i at $2i$ and $2i+1$
 - Parent at $\text{floor}(i/2)$
- Complete tree fills array with no gaps or overlap



Array-based heap

- Store node values in an array
- Root at index 1
- Children of index i at $2i$ and $2i+1$
 - Parent at $\text{floor}(i/2)$
- Complete tree fills array with no gaps or overlap
- Operations mostly the same
- **Min()**: return `arr[1]`
- **Insert(x)**:
 - Append to array
 - Percolate up
 - Increment size
 - Double capacity as needed
- **DeleteMin()**:
 - Swap `arr[1]` and `arr[size]`
 - Percolate `arr[1]` down
 - Decrement size and return `arr[size+1]`



Heapification

- Convert unsorted array into heap
 - Faster than multiple calls to insert
- **Algorithm:**
 - Call PercolateDown(i) from end to beginning
 - *Optimization:* don't percolate the leaves down
 - Avoids half of the heap
- **Complexity:** $\Theta(n)$ time
 - *Intuition:* half have no children, half of rest have 1 child, etc.
 - Only root has $\lg n$ levels below it
 - $\Theta(1)$ "on average"

```
1 Algorithm: Heapify(i)
2 for  $i = \lfloor n/2 \rfloor$  to 1 step  $-1$  do
3   | PercolateDown(i)
4 end
```

Priority queue implementations

Operation	Heap	Unsorted array	Sorted array	Balanced BST	Fibonacci heap
Insert(x)	$\Theta(\lg n)^*$	$\Theta(1)^*$	$\Theta(n)^*$	$O(\lg n)$	$\Theta(1)$
Max()	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
DeleteMax()	$\Theta(\lg n)$	$\Theta(n)$	$\Theta(1)$	$O(\lg n)$	$\Theta(\lg n)^*$
Build/Heapify	$\Theta(n)$	$\Theta(1)$	$\Theta(n \lg n)$	$\Theta(n \lg n)$	$\Theta(n)$

* amortized time

- BST has similar complexity, but higher coefficients
- Great at finding max (or min)
- Other operations (min/max, search, predecessor, etc.) are not good
 - Min-max heap can do either, but is more complex
- Fibonacci heap has even better complexity
 - More complex, higher coefficients, less space efficient
 - Fairly slow unless data is quite large

Heap exercise


1. Draw the max heap that is built from the array $\{8, 4, 6, 9, 2, 7, 1\}$.
2. Draw this heap at the end of every iteration of the **for** loop in HeapSort (below). You may ignore the effect of line 5.

Input: *data*: an array of integers to sort

Input: *n*: the number of values in data

Output: permutation of data such that
 $data[1] \leq \dots \leq data[n]$

```
1 Algorithm: HeapSort
2 data = MaxHeap.Build(data)
3 for i = n to 2 step -1 do
4   |   m = data.DeleteMax()
5   |   data[i] = m
6 end
7 return data
```




Heap sample solution


Tree view

Array view


Beginning:



After iter 1:




After iter 2:




Heap sample solution

Tree view


After iter 3:



After iter 4:



After iter 5:



After iter 6:



Array view



Union-Find data structure

- A.k.a., disjoint set data structure
- **Purpose:** represent partition of dataset
 - Identify whether elements belong to the same subset or not
- **Operations**
 - Initialize(n): set up each element ($1..n$) in its own subset
 - Find(x): return a partition ID for a given element
 - Union(x, y): combine subsets containing x and y together
- **Representation:** array of “pointers” (integers)
 - Find: follow pointers until you find a self-loop
 - Self-loop is partition ID (“root”)
 - Union: point Find(a) to b

Union-Find example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)


Union-Find example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

1	2	3	4	5	6	7
---	---	---	---	---	---	---

Pointers:




Union-Find example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	2	3	4	5	6	7
---	---	---	---	---	---	---

Pointers:




Union-Find example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	4	3	4	5	6	7
---	---	---	---	---	---	---

Pointers:




Union-Find example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	4	3	4	5	7	7
---	---	---	---	---	---	---

Pointers:




Union-Find example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	4	5	4	5	7	7
---	---	---	---	---	---	---

Pointers:




Union-Find example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:



2	4	5	4	6	7	7
---	---	---	---	---	---	---

Pointers:



Optimizing Union-Find

- Union complexity depends on Find
- Find complexity depends on height of tree
- Worst case: $\Theta(n)$
- First idea: add Find(a) to Find(b) (or vice versa)
- Second idea: add the smaller tree to the larger
- Third idea: flatten structure when we call Find()



Optimized example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)


Optimized example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

1	2	3	4	5	6	7
---	---	---	---	---	---	---

Pointers:




Optimized example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	2	3	4	5	6	7
---	---	---	---	---	---	---

Pointers:




Optimized example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	2	3	2	5	6	7
---	---	---	---	---	---	---

Pointers:




Optimized example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	2	3	2	5	7	7
---	---	---	---	---	---	---

Pointers:




Optimized example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	2	5	2	5	7	7
---	---	---	---	---	---	---

Pointers:




Optimized example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	2	5	2	7	7	7
---	---	---	---	---	---	---

Pointers:




Optimized example

- Show how the data structure changes after each of the following iterations:
 - Initialize(7)
 - Union(1, 2)
 - Union(1, 4)
 - Union(6, 7)
 - Union(3, 5)
 - Union(3, 6)
 - Find(3)

Array:

2	2	7	2	7	7	7
---	---	---	---	---	---	---

Pointers:



Union-Find operations

- **Find(x)**
 - Recursively point to answer
 - $\Theta(\alpha(n))$, amortized
 - Generally less than 4 for conceivable n
- **Union(a, b)**
 - Call Find on both sides first
 - Always point to larger tree
 - $\Theta(\alpha(n))$
- The Ackermann function
 - Incredibly fast-growing function
 - First few values:
 - 3, 7, 61, $2^{2^{2^{65536}}}$ – 3, ...

```
1 Algorithm: Find(x)
2 if unionfind[x]  $\neq$  x then
3   |   id = Find(unionfind[x]);
4   |   unionfind[x] = id;
5 end
6 return unionfind[x];
```


```
1 Algorithm: Union(a, b)
2 ra = Find(a);
3 rb = Find(b);
4 if size[ra] > size[rb] then
5   |   Swap ra and rb;
6 end
7 unionfind[ra] = rb;
8 size[rb] =
   size[ra] + size[rb];
```

Union-Find applications

- Major application: algorithm later in semester...
- Single linkage hierarchical clustering
 - All points start as separate, individual clusters (groups)
 - Repeatedly join "closest" clusters together
 - *Single linkage*: distance between clusters = min dist b/w points
- **Input:**
 - *data*: array of n points
 - c : number of clusters
- **Output:** c clusters
- **Pseudocode:**
 1. Initialize union-find
 2. Calculate all distances between points
 3. Repeat:
 4. If points with next smallest distance are not in same group:
 5. Union them
 6. Until there are c groups in union-find

Hierarchical clustering example

- **Pseudocode**
 - Union next closest points if not unioned
 - Repeat until desired # of clusters
- **Example:** *data* are integers, distance is abs. value, $c = 1$



$dist = 1$
 $dist = 2$
 $dist = 3$
 $dist = 4$

Union-Find exercise

- Perform single-linkage hierarchical clustering on the array below until everything is in one cluster
 - Distance is abs. value
- Draw the Union-Find data structure after each step



-4	-2	-1	3	4	7	9
----	----	----	---	---	---	---

- Rules
 - If given a choice, merge the cluster with the leftmost element
 - Merge elements left-to-right




Union-Find exercise

- Perform single-linkage hierarchical clustering on the array below until everything is in one cluster
 - Distance is abs. value
- Draw the Union-Find data structure after each step

-4	-2	-1	3	4	7	9
----	----	----	---	---	---	---

- Rules
 - If given a choice, merge the cluster with the leftmost element
 - Merge elements left-to-right




Union-Find exercise

- Perform single-linkage hierarchical clustering on the array below until everything is in one cluster
 - Distance is abs. value
- Draw the Union-Find data structure after each step

-4	-2	-1	3	4	7	9
----	----	----	---	---	---	---

- Rules
 - If given a choice, merge the cluster with the leftmost element
 - Merge elements left-to-right




Union-Find exercise

- Perform single-linkage hierarchical clustering on the array below until everything is in one cluster
 - Distance is abs. value
- Draw the Union-Find data structure after each step

-4	-2	-1	3	4	7	9
----	----	----	---	---	---	---

- Rules
 - If given a choice, merge the cluster with the leftmost element
 - Merge elements left-to-right

Union(3, 4)
or: Union(-1, 3)



Array representation (final):

3	3	7	7	7	7	7
---	---	---	---	---	---	---

Note that array values represent the *position* (index) of the element, not its value.

For example, element #1 (-4) points to element #3 (-1), which points to element #7 (9).

Coming up

- Union-Find
- Sorting
- Search
- **Recommended reading:** Sections 5.1, 5.3, 5.4 (just "Bottom-Up Heap Construction" to the end), 7.1, and 7.3
 - *Practice problems:* R-5.8, R-5.9, R-5.10, R-5.12, C-5.5, A-5.1, R-7.7, R-7.8, C-7.8, A-7.1